

- ✓ Import necessary library

```
1 import pandas as pd
2 import numpy as np
3 import torch
4 from torchvision import datasets
5 from torchvision.transforms import ToTensor
6 import torchvision.transforms as transforms
7 from torch.utils.data import DataLoader
8 import matplotlib.pyplot as plt
9 import torch.nn as nn
10 import torch.nn.functional as F
11 from torch import optim
12
13 from sklearn.metrics import confusion_matrix, recall_score, precision_score, f1_score, accuracy_score
14 from sklearn.metrics import classification_report, ConfusionMatrixDisplay
```

- Dataset preparation

- Acquire the MNIST dataset
- Preprocess the data (normalizaon, reshaping, train/test split)

About Dataset ref: <https://www.kaggle.com/datasets/oddrational/mnist-in-csv>

MNIST contains a collection of 70,000, 28 x 28 images of handwritten digits from 0 to 9.

The dataset consists of two files:

- train : 60,000 training examples and labels
- test: 10,000 test examples and labels

each set - label: a number from 0 to 9,image: the pixel values (a number from 0 to 255).

```
1 train_data = datasets.MNIST(
2     root = 'data',
3     train = True,
4     transform = ToTensor(),
5     download = True,
6 )
7 test_data = datasets.MNIST(
8     root = 'data',
9     train = False,
10    transform = ToTensor(),
11    download = True,
12 )
```

1 test data

```
Dataset MNIST
  Number of datapoints: 10000
  Root location: data
  Split: Test
  StandardTransform
Transform: ToTensor()
```

```
1 train data.data[0]
```

```
tensor([[ 0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,
         0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0],
        [ 0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,
         0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0],
        [ 0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,
         0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0],
        [ 0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,
         0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0],
        [ 0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,
         0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0],
        [ 0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,
         0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0],
        [ 0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,
         0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0],
        [ 0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,   0,
         0,   0,   0,   0,   0,   0,   0,   0,   3,  18,
        18,  18, 126, 136, 175,  26, 166, 255, 247, 127,   0,   0,   0,   0]
```

```
[ 0, 0, 0, 0, 0, 0, 0, 0, 30, 36, 94, 154, 170, 253,
 253, 253, 253, 253, 225, 172, 253, 242, 195, 64, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 49, 238, 253, 253, 253, 253,
 253, 253, 253, 251, 93, 82, 82, 56, 39, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 18, 219, 253, 253, 253, 253,
 198, 182, 247, 241, 0, 0, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 80, 156, 107, 253, 253, 205,
 11, 0, 43, 154, 0, 0, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 14, 1, 154, 253, 90,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 139, 253, 190,
 2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 11, 190, 253,
 70, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 35, 241,
 225, 160, 108, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 81,
 240, 253, 253, 119, 25, 0, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 45, 186, 253, 253, 150, 27, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 16, 93, 252, 253, 187, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 249, 253, 249, 64, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 46, 130, 183, 253, 253, 207, 2, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 39, 148,
 229, 253, 253, 253, 250, 182, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 24, 114, 221, 253,
 253, 253, 253, 201, 78, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 23, 66, 213, 253, 253, 253,
 253, 198, 81, 2, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 18, 171, 219, 253, 253, 253, 253, 195,
 80, 9, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 55, 172, 226, 253, 253, 253, 253, 244, 133, 11,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 136, 253, 253, 253, 212, 135, 132, 16, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0],
[ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
 0, 0, 0, 0, 0, 0, 0, 0]]],
dtype=torch.uint8)
```

▼ Normalization

```
1 #Get mean and std of training data
2
3 print('Min Pixel Value: {} \nMax Pixel Value: {}'.format(train_data.data.min(), train_data.data.max()))
4 print('Mean Pixel Value {} \nPixel Values Std: {}'.format(train_data.data.float().mean(), train_data.data.float().std()))
5 print('Scaled Mean Pixel Value {} \nScaled Pixel Values Std: {}'.format(train_data.data.float().mean(), train_data.data.float().std()))
```

```
Min Pixel Value: 0
Max Pixel Value: 255
Mean Pixel Value 33.31842041015625
Pixel Values Std: 78.56748962402344
Scaled Mean Pixel Value 0.13066047430038452
Scaled Pixel Values Std: 0.30810779333114624
```

```
1 #Get mean and std of testing data
2
3 print('Min Pixel Value: {} \nMax Pixel Value: {}'.format(test_data.data.min(), test_data.data.max()))
4 print('Mean Pixel Value {} \nPixel Values Std: {}'.format(test_data.data.float().mean(), test_data.data.float().std()))
5 print('Scaled Mean Pixel Value {} \nScaled Pixel Values Std: {}'.format(test_data.data.float().mean(), test_data.data.float().std()))
```

```
Min Pixel Value: 0
Max Pixel Value: 255
Mean Pixel Value 33.79122543334961
Pixel Values Std: 79.17247009277344
Scaled Mean Pixel Value 0.1325146108865738
Scaled Pixel Values Std: 0.3104802668094635
```

```
1 transform_train=transforms.Compose([
2     transforms.ToTensor(),
3     transforms.Normalize((0.1307,), (0.3081,)) # divides by 255
4 ])
5 transform_test=transforms.Compose([
6     transforms.ToTensor(),
7     transforms.Normalize((0.1307,), (0.3081,)) # divides by 255
8 ])
```

1

```
1 test data.data[1]
```

3/24

```
dtype=torch.uint8)
```

```
2 train_loader = torch.utils.data.DataLoader(train_data,
3                                             batch_size=batch_size,
4                                             shuffle=True),
5 'test' : torch.utils.data.DataLoader(test_data,
6                                       batch_size=batch_size,
7                                       shuffle=True)}
8 loaders
```

```
{'train': <torch.utils.data.dataloader.DataLoader at 0x7fc27e4fb2b0>,
 'test': <torch.utils.data.dataloader.DataLoader at 0x7fc27e4fb250>}
```

✓ Show sample data

```
1 examples = enumerate(loaders['test'])
2 batch_id, (example_data, example_targets) = next(examples)
```

```
1 example_data.shape
   torch.Size([100, 1, 28, 28])
```

```
1 fig = plt.figure()
2
3 for i in range(6):
4     plt.subplot(2,3,i+1)
5     plt.tight_layout()
6     plt.imshow(example_data[i][0], cmap='gray', interpolation='none')
7     plt.title("Ground Truth: {}".format(example_targets[i]))
8     plt.xticks([])
9     plt.yticks([])
10 fig
```

✓ CNN model design

OUTPUT FORMULA FOR CONVOLUTION

$$\bullet O = \frac{W-K+2P}{S} + 1$$

OUTPUT FORMULA FOR POOLING

$$\bullet O = \frac{W-K}{S} + 1$$

- O : output height/length
- W : input height/length
- K : filter size (kernel size)
- P : padding
 - $P = \frac{K-1}{2}$
- S : stride

Ref: <https://www.run.ai/guides/deep-learning-for-computer-vision/pytorch-cnn>

ดับเบิลคลิก (หรือกด Enter) เพื่อแก้ไข

✓ Def for calculate next output

```
1 def cal_out(input,k,s,p):
2     out = ((input-k+(2*p))/s)+1
3     return out
```

✓ Model

```
1 class CNN(nn.Module):
2     def __init__(self):
3         super(CNN, self).__init__()
4         self.conv1 = nn.Conv2d(in_channels=1,out_channels=100,kernel_size=6,stride=2,padding=1)
5         self.relu1 = nn.ReLU()
6         self.pool1 = nn.MaxPool2d(kernel_size=4,stride=2)
7         self.out1 = nn.Linear(100 * 6 * 6, 10)
8         self.relu3 = nn.ReLU()
9
10    def forward(self, x):
11        x = self.conv1(x)
12        x = self.relu1(x)
13        x = self.pool1(x)
14
15        # flatten the output of conv2
16        x = x.view(x.size(0), -1)
17        output = self.out1(x)
18        return output
```

```

1 # class CNN(nn.Module):
2 #     def __init__(self):
3 #         super(CNN, self).__init__()
4 #         self.conv1 = nn.Conv2d(in_channels=1,out_channels=4,kernel_size=7,stride=1,padding=
5 #         self.relu1 = nn.ReLU()
6 #         self.pool1 = nn.MaxPool2d(kernel_size=3)
7 #         self.conv2 = nn.Conv2d(in_channels=4,out_channels=16,kernel_size=2,stride=1,padding=
8 #         self.relu2 = nn.ReLU()
9 #         self.out1 = nn.Linear(16 * 9 * 9, 10)
10 #         self.relu3 = nn.ReLU()
11
12 #     def forward(self, x):
13 #         x = self.conv1(x)
14 #         x = self.relu1(x)
15 #         x = self.pool1(x)
16 #         x = self.conv2(x)
17 #         x = self.relu2(x)
18 #         # flatten the output of conv2
19 #         x = x.view(x.size(0), -1)
20 #         output = self.out1(x)
21 #         return output

```

```

1 model = CNN()
2 print(model)

```

```

CNN(
  (conv1): Conv2d(1, 100, kernel_size=(6, 6), stride=(2, 2), padding=(2, 2))
  (relu1): ReLU()
  (pool1): MaxPool2d(kernel_size=4, stride=2, padding=0, dilation=1, ceil_mode=False)
  (out1): Linear(in_features=3600, out_features=10, bias=True)
  (relu3): ReLU()
)

```

✓ Define loss function, optimization algorithm

```

1 def call_fuction(model):
2     loss_func = nn.CrossEntropyLoss()
3     optimizer = optim.Adam(model.parameters(), lr = learning_rate, weight_decay=0.1)
4     return loss_func,optimizer

```

```

1 loss_func,optimizer = call_fuction(model)

```

```

1 optimizer

```

```

Adam (
  Parameter Group 0
    amsgrad: False
    betas: (0.9, 0.999)
    capturable: False
    differentiable: False
    eps: 1e-08
    foreach: None
    fused: None
    lr: 0.01
    maximize: False
    weight_decay: 0.1
)

```

✓ Train model

```

1 print('Number of iterations in each epoch in train set =',(len(train_data) / batch_size))
2 print('Number of iterations in each epoch in test set =',(len(test_data) / batch_size))
3 print('Number of epochs =',num_epochs)

```

```

Number of iterations in each epoch in train set = 600.0
Number of iterations in each epoch in test set = 100.0
Number of epochs = 8

```

```

1 def train(model,num_epochs,optimizer,loss_func):
2     iter = 0
3     correctImages = 0
4     totalImages = 0
5     trainingLoss = 0
6     model.train()
7     loss_each = []
8     acc_each = []
9     for epoch in range(num_epochs):
10
11         for i, (images, labels) in enumerate(loaders['train']):
12             images = images.requires_grad_()
13             optimizer.zero_grad()
14             outputs = model(images)
15             _, predicted = torch.max(outputs, 1)
16             loss = loss_func(outputs, labels)
17             loss.backward()
18             optimizer.step()
19             iter += 1
20
21             correctImages += (predicted == labels).sum().item()
22             totalImages += labels.size(0)
23             trainingLoss += loss.item()
24             accumulateAccuracy = round((correctImages/totalImages)*100,4)
25
26             # Number of iterations in each epoch in train set
27             if iter % 600 == 0:
28                 print('Epoch [{}/{}] Iteration: {}'.format(epoch+1, num_epochs, iter))
29                 print('Training accuracy: {} loss: {}'.format(accumulateAccuracy, round(loss.
30                     loss_each.append(round(loss.item(),4))
31                     acc_each.append(accumulateAccuracy)
32                     pass
33             pass
34         pass
35     return loss_each,acc_each

```

```
1 trainloss,trainacc = train(model,num_epochs,optimizer,loss_func)
```

```
Epoch [1/8] Iteration: 600
Training accuracy: 89.1883 loss: 0.2394
```

```
KeyboardInterrupt                                Traceback (most recent call last)
<ipython-input-23-17480899683a> in <cell line: 1>()
----> 1 trainloss,trainacc = train(model,num_epochs,optimizer,loss_func)
```

```

^ 6 frames
/usr/local/lib/python3.10/dist-packages/torch/nn/functional.py in _max_pool2d(input, kernel_size, stride, padding, dilat:
780     if stride is None:
781         stride = torch.jit.annotate(List[int], [])
--> 782     return torch.max_pool2d(input, kernel_size, stride, padding, dilation, ceil_mode)
783
784

```

```
KeyboardInterrupt:
```

✓ Model evaluation


```

1 def test(cnn):
2     # Test the model
3     cnn.eval()
4     loss_each = []
5     acc_each = []
6     y_trues = []
7     y_preds = []
8     with torch.no_grad():
9         iter = 0
10        correctImages = 0
11        totalImages = 0
12
13        for epoch in range(num_epochs):
14            for images, labels in loaders['test']:
15                outputs = cnn(images)
16                _, predicted = torch.max(outputs, 1)
17                loss = loss_func(outputs, labels)
18                if epoch == (num_epochs-1):
19                    for i in labels.tolist():
20                        y_trues.append(i)
21                    for i in predicted.tolist():
22                        y_preds.append(i)
23                correctImages += (predicted == labels).sum().item()
24                totalImages += labels.size(0)
25                accumulateAccuracy = round((correctImages/totalImages)*100,4)
26                iter += 1
27
28            # Number of iterations in each epoch in test set
29            if iter % 100 == 0:
30                print('Epoch [{}/{}] Iteration: {}'.format(epoch+1, num_epochs, iter))
31                print('Testing accuracy: {} loss: {}'.format(accumulateAccuracy, round(loss.item(),4)))
32                loss_each.append(round(loss.item(),4))
33                acc_each.append(accumulateAccuracy)
34                pass
35            pass
36        pass
37    pass
38
39    # pass
40    return y_trues, y_preds, loss_each, acc_each

1 y_trues, y_preds, testloss, testacc = test(model)

```

✓ Plot Loss and Accuracy

```

1 def plot_acc():
2     epochs = range(1,num_epochs+1)
3     plt.plot(epochs, trainacc, 'ro', label='Accuracy of Train set')
4     plt.plot(epochs, testacc, 'bo', label='Accuracy of Test set')
5     plt.title('Train Vs Test accuracy')
6     plt.xlabel('Epoch')
7     plt.ylabel('Accuracy')
8     plt.legend(loc=0)
9     plt.figure()
10    plt.show()

1 plot_acc()

1 def plot_loss():
2     epochs = range(1,num_epochs+1)
3     plt.plot(epochs, trainloss, 'r', label='Loss of Train set')
4     plt.plot(epochs, testloss, 'b--', label='Loss of Test set')
5     plt.title('Train Vs Test Loss')
6     plt.xlabel('Epoch')
7     plt.ylabel('Loss')
8     plt.legend(loc=0)

```

```

8     plt.legend(loc=0,
9     plt.figure()
10    plt.show()

```

```
1 plot_loss()
```

✓ Confusion Matrix

```
1 list(range(0,10))
```

```

1 def disp_report_and_cm():
2     print(classification_report(y_trues, y_preds))
3     cm = confusion_matrix(y_trues, y_preds)
4     disp = ConfusionMatrixDisplay(confusion_matrix=cm,
5                                   display_labels=list(range(0,10)))
6     disp.plot()
7     plt.show()

```

```
1 disp_report_and_cm()
```

✓ Result before Optimize

```

1 print('Accuracy:',accuracy_score(y_trues, y_preds))
2 print('Precision:',precision_score(y_trues, y_preds, average='macro'))
3 print('Recall:',recall_score(y_trues, y_preds, average='macro'))
4 print('f1-Score:',f1_score(y_trues, y_preds, average='macro'))

1 def plot_predict(model):
2     with torch.no_grad():
3         output = model(example_data)
4         fig = plt.figure()
5         for i in range(6):
6             plt.subplot(2,3,i+1)
7             plt.tight_layout()
8             plt.imshow(example_data[i][0], cmap='gray', interpolation='none')
9             plt.title("Prediction: {}".format(
10                 output.data.max(1, keepdim=True)[1][i].item()))
11             plt.xticks([])
12             plt.yticks([])
13     return(fig)

```

```
1 plot_predict(model)
```

✓ Model optimization

For repeatable experiments we have to set random seeds for anything using random number generation - this means numpy and random as well! It's also worth mentioning that cuDNN uses nondeterministic algorithms which can be disabled setting `torch.backends.cudnn.enabled = False`. (การสุ่มเพื่อเทรน cnn แต่ละครั้งจะ nondeterministic หรือการที่มีความคล้ายกันแต่ไม่เหมือนกัน)

```

1 def before_repeat():
2     random_seed = 1
3     torch.backends.cudnn.enabled = False
4     torch.manual_seed(random_seed)

```

```
1 before_repeat()
```

ก่อน train model ใหม่ต้องลบรันทาม และรัน def call_function(),train(),test() ก่อนด้วย

✓ Different CNN architectures

เพิ่ม drop out, fully-connected, ลด padding

```

1 # cal_out(7,5,1,2)

1 batch_size = 100
2 num_workers = 6
3 n_iters = 10000
4
5 #จำนวน epoch = จำนวน iteration/จำนวน batch (training set = 60000)
6 num_epochs = n_iters / (len(train_data) / batch_size)
7 num_epochs = int(num_epochs)
8 print(num_epochs)
9
10 learning_rate = 0.001

16

1 class CNN1(nn.Module):
2     def __init__(self):
3         super(CNN1, self).__init__()
4         self.conv1 = nn.Conv2d(in_channels=1, out_channels=16, kernel_size=5, stride=1, padding=0)
5         self.relu1 = nn.ReLU()
6         self.pool1 = nn.MaxPool2d(kernel_size=2)
7         self.conv2 = nn.Conv2d(16, 32, 5, 1, 0)
8         self.relu2 = nn.ReLU()
9         self.pool2 = nn.MaxPool2d(kernel_size=2)
10        self.dropout = nn.Dropout(p=0.3)
11
12        self.out1 = nn.Linear(32 * 4 * 4, 256)
13        # self.out1 = nn.Linear(32 * 4 * 4, 10)
14        self.relu3 = nn.ReLU()
15
16        #fully connected layer, output 10 classes
17        self.out2 = nn.Linear(256, 10)
18
19
20    def forward(self, x):
21        x = self.conv1(x)
22        x = self.relu1(x)
23        x = self.pool1(x)
24
25        x = self.conv2(x)
26        x = self.relu2(x)
27        x = self.pool2(x)
28
29        x = self.dropout(x)
30
31        # # flatten the output of conv2
32        x = x.view(x.size(0), -1)
33        x = self.out1(x)
34        x = self.relu3(x)
35        x = self.dropout(x)
36        output = self.out2(x)
37        return output

1 model1 = CNN1()
2 print(model1)

```

```

CNN1(
  (conv1): Conv2d(1, 16, kernel_size=(5, 5), stride=(1, 1))
  (relu1): ReLU()
  (pool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): Conv2d(16, 32, kernel_size=(5, 5), stride=(1, 1))
  (relu2): ReLU()
  (pool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (dropout): Dropout(p=0.3, inplace=False)
  (out1): Linear(in_features=512, out_features=256, bias=True)

```

```

    (relu3): ReLU()
    (out2): Linear(in_features=256, out_features=10, bias=True)
)

1 loss_func,optimizer = call_fuction(model1)
2 optimizer

Adam (
Parameter Group 0
  amsgrad: False
  betas: (0.9, 0.999)
  capturable: False
  differentiable: False
  eps: 1e-08
  foreach: None
  fused: None
  lr: 0.001
  maximize: False
  weight_decay: 0.1
)

1 trainloss,trainacc = train(model1,num_epochs,optimizer,loss_func)

Epoch [1/16] Iteration: 600
Training accuracy: 87.0333 loss: 0.2743
Epoch [2/16] Iteration: 1200
Training accuracy: 88.9292 loss: 0.2899
Epoch [3/16] Iteration: 1800
Training accuracy: 89.5906 loss: 0.2048
Epoch [4/16] Iteration: 2400
Training accuracy: 89.9567 loss: 0.3229
Epoch [5/16] Iteration: 3000
Training accuracy: 90.1717 loss: 0.3549
Epoch [6/16] Iteration: 3600
Training accuracy: 90.3083 loss: 0.3259
Epoch [7/16] Iteration: 4200
Training accuracy: 90.3957 loss: 0.3827
Epoch [8/16] Iteration: 4800
Training accuracy: 90.475 loss: 0.4357
Epoch [9/16] Iteration: 5400
Training accuracy: 90.5202 loss: 0.4036
Epoch [10/16] Iteration: 6000
Training accuracy: 90.589 loss: 0.4129
Epoch [11/16] Iteration: 6600
Training accuracy: 90.6445 loss: 0.3454
Epoch [12/16] Iteration: 7200
Training accuracy: 90.6831 loss: 0.3359
Epoch [13/16] Iteration: 7800
Training accuracy: 90.7206 loss: 0.3274
Epoch [14/16] Iteration: 8400
Training accuracy: 90.7415 loss: 0.3308
Epoch [15/16] Iteration: 9000
Training accuracy: 90.7627 loss: 0.2387
Epoch [16/16] Iteration: 9600
Training accuracy: 90.7754 loss: 0.2857


1 y_trues, y_preds, testloss, testacc = test(model1)

Epoch [1/16] Iteration: 100
Testing accuracy: 93.65 loss: 0.2445
Epoch [2/16] Iteration: 200
Testing accuracy: 93.65 loss: 0.2651
Epoch [3/16] Iteration: 300
Testing accuracy: 93.65 loss: 0.172
Epoch [4/16] Iteration: 400
Testing accuracy: 93.65 loss: 0.2816
Epoch [5/16] Iteration: 500
Testing accuracy: 93.65 loss: 0.3953
Epoch [6/16] Iteration: 600
Testing accuracy: 93.65 loss: 0.2993
Epoch [7/16] Iteration: 700
Testing accuracy: 93.65 loss: 0.2466
Epoch [8/16] Iteration: 800
Testing accuracy: 93.65 loss: 0.1761
Epoch [9/16] Iteration: 900
Testing accuracy: 93.65 loss: 0.2351
Epoch [10/16] Iteration: 1000
Testing accuracy: 93.65 loss: 0.2252
Epoch [11/16] Iteration: 1100
Testing accuracy: 93.65 loss: 0.2347
Epoch [12/16] Iteration: 1200
Testing accuracy: 93.65 loss: 0.3088
Epoch [13/16] Iteration: 1300
Testing accuracy: 93.65 loss: 0.2522
Epoch [14/16] Iteration: 1400
Testing accuracy: 93.65 loss: 0.2311

```

```
Epoch [15/16] Iteration: 1500  
Testing accuracy: 93.65 loss: 0.2822  
Epoch [16/16] Iteration: 1600  
Testing accuracy: 93.65 loss: 0.2294
```

```
1 plot_acc()
```




```
1 plot_loss()
```



```
1 disp_report_and_cm()
```

	precision	recall	f1-score	support
0	0.97	0.96	0.96	980
1	0.96	0.99	0.98	1135
2	0.94	0.89	0.92	1032
3	0.88	0.97	0.93	1010
4	0.92	0.97	0.95	982
5	0.97	0.94	0.96	892
6	0.91	0.97	0.94	958
7	0.98	0.87	0.92	1028
8	0.92	0.91	0.92	974
9	0.91	0.90	0.90	1009
accuracy			0.94	10000
macro avg	0.94	0.94	0.94	10000
weighted avg	0.94	0.94	0.94	10000



```

1 print('Accuracy:',accuracy_score(y_trues, y_preds))
2 print('Precision:',precision_score(y_trues, y_preds, average='macro'))
3 print('Recall:',recall_score(y_trues, y_preds, average='macro'))
4 print('f1-Score:',f1_score(y_trues, y_preds, average='macro'))


```

```

Accuracy: 0.9365
Precision: 0.9376093337555854
Recall: 0.9364077468081048
f1-Score: 0.9361741060721295

```

```
1 plot_predict(model1)
```



✓ Data augmentation

```

1 transform_train=transforms.Compose([
2     transforms.RandomHorizontalFlip(),
3     transforms.RandomVerticalFlip(),
4     transforms.ToTensor(),
5     transforms.Normalize((0.1307,), (0.3081,)) # divides by 255
6 ])
7 transform_test=transforms.Compose([
8     transforms.ToTensor(),
9     transforms.Normalize((0.1325,), (0.3105,)) # divides by 255
10 ])

```

```

1 train_data = datasets.MNIST(
2     root = 'data',
3     train = True,
4     transform = transform_train,
5     download = True,
6 )
7 test_data = datasets.MNIST(
8     root = 'data',
9     train = False,
10    transform = transform_test,
11    download = True,
12 )

```

```

1 batch_size = 100
2 num_workers = 6
3 n_iters = 10000
4
5 #จำนวน epoch = จำนวน iteration/จำนวน batch (training set = 60000)
6 num_epochs = n_iters / (len(train_data) / batch_size)
7 num_epochs = int(num_epochs)
8 print(num_epochs)
9
10 learning_rate = 0.001

```

16

```

1 loaders = {
2     'train' : torch.utils.data.DataLoader(train_data,
3                                           batch_size=batch_size,
4                                           shuffle=True),
5     'test'  : torch.utils.data.DataLoader(test_data,
6                                           batch_size=batch_size,
7                                           shuffle=True)}
8 loaders

```

```

{'train': <torch.utils.data.dataloader.DataLoader at 0x7fdb7e23fe80>,
 'test': <torch.utils.data.dataloader.DataLoader at 0x7fdb7e23ef20>}

```

```

1 examples = enumerate(loaders['test'])
2 batch_id, (example_data, example_targets) = next(examples)

```

```


1 example_data.shape
torch.Size([100, 1, 28, 28])

```

```

1 fig = plt.figure()
2
3 for i in range(6):
4     plt.subplot(2,3,i+1)
5     plt.tight_layout()
6     plt.imshow(example_data[i][0], cmap='gray', interpolation='none')
7     plt.title("Ground Truth: {}".format(example_targets[i]))
8     plt.xticks([])
9     plt.yticks([])
10 fig

```



```

1 loss_func,optimizer = call_fuction(model1)
2 optimizer

```

```

Adam (
Parameter Group 0
  amsgrad: False
  betas: (0.9, 0.999)
  capturable: False
  differentiable: False
  eps: 1e-08
  foreach: None
  fused: None
  lr: 0.001
  maximize: False
  weight_decay: 0.1
)

```

```

1 trainloss,trainacc = train(model1,num_epochs,optimizer,loss_func)

```

```

Epoch [1/16] Iteration: 600
Training accuracy: 68.955 loss: 0.6889
Epoch [2/16] Iteration: 1200
Training accuracy: 71.1675 loss: 0.6685
Epoch [3/16] Iteration: 1800

```




```
Training accuracy: 72.2644 loss: 0.8567
Epoch [4/16] Iteration: 2400
Training accuracy: 73.0175 loss: 0.7257
Epoch [5/16] Iteration: 3000
Training accuracy: 73.4433 loss: 0.6579
Epoch [6/16] Iteration: 3600
Training accuracy: 73.8172 loss: 0.7397
Epoch [7/16] Iteration: 4200
Training accuracy: 74.061 loss: 0.9463
Epoch [8/16] Iteration: 4800
Training accuracy: 74.2517 loss: 0.9383
Epoch [9/16] Iteration: 5400
Training accuracy: 74.3863 loss: 0.9362
Epoch [10/16] Iteration: 6000
Training accuracy: 74.542 loss: 0.6689
Epoch [11/16] Iteration: 6600
Training accuracy: 74.6724 loss: 0.7209
Epoch [12/16] Iteration: 7200
Training accuracy: 74.7761 loss: 0.7705
Epoch [13/16] Iteration: 7800
Training accuracy: 74.884 loss: 0.8183
Epoch [14/16] Iteration: 8400
Training accuracy: 74.9385 loss: 0.7529
Epoch [15/16] Iteration: 9000
Training accuracy: 75.0258 loss: 0.6349
Epoch [16/16] Iteration: 9600
Training accuracy: 75.0896 loss: 0.9992
```

```
1 y_trues, y_preds, testloss, testacc = test(model1)
```

```
Epoch [1/16] Iteration: 100
Testing accuracy: 82.4 loss: 0.6508
Epoch [2/16] Iteration: 200
Testing accuracy: 82.4 loss: 0.713
Epoch [3/16] Iteration: 300
Testing accuracy: 82.4 loss: 0.6943
Epoch [4/16] Iteration: 400
Testing accuracy: 82.4 loss: 0.7386
Epoch [5/16] Iteration: 500
Testing accuracy: 82.4 loss: 0.5758
Epoch [6/16] Iteration: 600
Testing accuracy: 82.4 loss: 0.6414
Epoch [7/16] Iteration: 700
Testing accuracy: 82.4 loss: 0.5768
Epoch [8/16] Iteration: 800
Testing accuracy: 82.4 loss: 0.6378
Epoch [9/16] Iteration: 900
Testing accuracy: 82.4 loss: 0.6242
Epoch [10/16] Iteration: 1000
Testing accuracy: 82.4 loss: 0.5927
Epoch [11/16] Iteration: 1100
Testing accuracy: 82.4 loss: 0.6709
Epoch [12/16] Iteration: 1200
Testing accuracy: 82.4 loss: 0.5382
Epoch [13/16] Iteration: 1300
Testing accuracy: 82.4 loss: 0.672
Epoch [14/16] Iteration: 1400
Testing accuracy: 82.4 loss: 0.7175
Epoch [15/16] Iteration: 1500
Testing accuracy: 82.4 loss: 0.6818
Epoch [16/16] Iteration: 1600
Testing accuracy: 82.4 loss: 0.6097
```

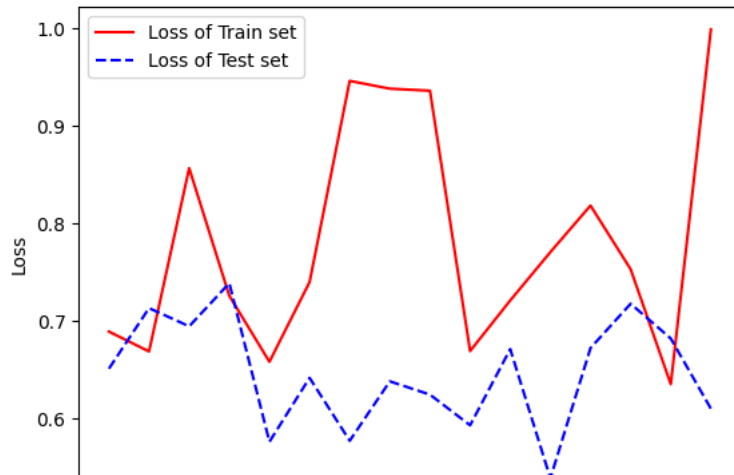
```
1 plot_acc()
```

Train Vs Test accuracy




```
1 plot_loss()
```

Train Vs Test Loss



```
1 disp_report_and_cm()
```


	precision	recall	f1-score	support
0	0.94	0.93	0.93	980
1	0.97	0.98	0.97	1135
2	0.80	0.72	0.75	1032
3	0.62	0.90	0.74	1010
4	0.93	0.92	0.92	982
5	0.72	0.68	0.70	892
6	0.90	0.75	0.82	958
7	0.80	0.89	0.84	1028
8	0.98	0.58	0.73	974
9	0.74	0.86	0.80	1009
accuracy			0.82	10000
macro avg	0.84	0.82	0.82	10000
weighted avg	0.84	0.82	0.82	10000



```
1 print('Accuracy:',accuracy_score(y_trues, y_preds))
2 print('Precision:',precision_score(y_trues, y_preds, average='macro'))
3 print('Recall:',recall_score(y_trues, y_preds, average='macro'))
4 print('f1-Score:',f1_score(y_trues, y_preds, average='macro'))
```

Accuracy: 0.824
 Precision: 0.839832031379148
 Recall: 0.8198635829151002
 f1-Score: 0.82055084290603

```
1 plot_predict(model1)
```



✓ Test by testing dataset (50 images)

```
1 from google.colab import drive
2 drive.mount('/content/drive')

1 transform_test=transforms.Compose([
2     transforms.Grayscale(),
3     transforms.ToTensor(),
4     transforms.Normalize((0.1325,), (0.3105,)) # divides by 255
5 ])

1 from torchvision.datasets import ImageFolder
2 test_data_new = ImageFolder('/content/drive/MyDrive/Senior_Project/50images/testing/', transform=transform_test)

1 print('New test images :', len(test_data_new))

1 loaders['test_new'] = torch.utils.data.DataLoader(test_data_new,
2                                                     batch_size=batch_size,
3                                                     shuffle=True)
```


```

1 examples = enumerate(loaders['test_new'])
2 batch_id, (example_data, example_targets) = next(examples)

1 example_data.shape
   torch.Size([100, 1, 28, 28])

1 fig = plt.figure()
2
3 for i in range(6):
4     plt.subplot(2,3,i+1)
5     plt.tight_layout()
6     plt.imshow(example_data[i][0], cmap='gray', interpolation='none')
7     plt.title("Ground Truth: {}".format(example_targets[i]))
8     plt.xticks([])
9     plt.yticks([])
10 fig

```



```

1 print('Number of iterations in each epoch in test set =',(len(test_data_new) / batch_size))

   Number of iterations in each epoch in test set = 5.0

```

```

1 def test2(cnn):
2     # Test the model
3     cnn.eval()
4     loss_each = []
5     acc_each = []
6     y_trues = []
7     y_preds = []
8     with torch.no_grad():
9         iter = 0
10        correctImages = 0
11        totalImages = 0
12        for epoch in range(num_epochs):
13            for images, labels in loaders['test_new']:
14                outputs = cnn(images)
15                _, predicted = torch.max(outputs, 1)
16                loss = loss_func(outputs, labels)
17                if epoch == (num_epochs-1):
18                    for i in labels.tolist():
19                        y_trues.append(i)
20                    for i in predicted.tolist():
21                        y_preds.append(i)
22                correctImages += (predicted == labels).sum().item()
23                totalImages += labels.size(0)
24                accumulateAccuracy = round((correctImages/totalImages)*100,4)
25                iter += 1
26
27            # Number of iterations in each epoch in test set
28            if iter % 5 == 0:
29                print('Epoch [{}]/[{}] Iteration: {}'.format(epoch+1, num_epochs, iter))
30                print('Testing accuracy: {} loss: {}'.format(accumulateAccuracy, round(loss.item(),4)))
31                loss_each.append(round(loss.item(),4))
32                acc_each.append(accumulateAccuracy)
33                pass
34            pass
35        pass
36    pass
37        # print('Test Accuracy of the model on the 10000 test images: %.2f' % accuracy)
38
39    # pass
40    return y_trues, y_preds, loss_each, acc_each

```

✓ Base Model

```

1 y_trues, y_preds, testloss, testacc = test2(model)

1 plot_acc()

1 plot_loss()

1 disp_report_and_cm()

1 print('Accuracy:', accuracy_score(y_trues, y_preds))
2 print('Precision:', precision_score(y_trues, y_preds, average='macro'))
3 print('Recall:', recall_score(y_trues, y_preds, average='macro'))
4 print('f1-Score:', f1_score(y_trues, y_preds, average='macro'))

1 plot_predict(model)

```

› New Model


[] ↳ ขอน 6 เซลล์

✓ New Model with augmentation

```
1 y_trues, y_preds, testloss, testacc = test2(model1)
```

```
Epoch [1/16] Iteration: 5  
Testing accuracy: 42.8 loss: 1.8082  
Epoch [2/16] Iteration: 10  
Testing accuracy: 42.8 loss: 1.6302  
Epoch [3/16] Iteration: 15  
Testing accuracy: 42.8 loss: 1.7968  
Epoch [4/16] Iteration: 20  
Testing accuracy: 42.8 loss: 1.7691  
Epoch [5/16] Iteration: 25  
Testing accuracy: 42.8 loss: 1.6936  
Epoch [6/16] Iteration: 30  
Testing accuracy: 42.8 loss: 1.6883  
Epoch [7/16] Iteration: 35  
Testing accuracy: 42.8 loss: 1.6485  
Epoch [8/16] Iteration: 40  
Testing accuracy: 42.8 loss: 1.8405  
Epoch [9/16] Iteration: 45  
Testing accuracy: 42.8 loss: 1.5212  
Epoch [10/16] Iteration: 50  
Testing accuracy: 42.8 loss: 1.8063  
Epoch [11/16] Iteration: 55  
Testing accuracy: 42.8 loss: 1.7826  
Epoch [12/16] Iteration: 60  
Testing accuracy: 42.8 loss: 1.8067  
Epoch [13/16] Iteration: 65  
Testing accuracy: 42.8 loss: 1.6807  
Epoch [14/16] Iteration: 70  
Testing accuracy: 42.8 loss: 1.6512  
Epoch [15/16] Iteration: 75  
Testing accuracy: 42.8 loss: 1.757  
Epoch [16/16] Iteration: 80  
Testing accuracy: 42.8 loss: 1.6635
```

```
1 plot_acc()
```




```
1 plot_loss()
```



```
1 disp_report_and_cm()
```

	precision	recall	f1-score	support
0	0.76	0.78	0.77	50
1	0.36	0.34	0.35	50
2	0.62	0.42	0.50	50
3	0.40	0.48	0.44	50
4	0.26	0.60	0.36	50
5	1.00	0.02	0.04	50
6	0.61	0.44	0.51	50
7	0.94	0.30	0.45	50
8	0.31	0.54	0.39	50
9	0.36	0.36	0.36	50
accuracy			0.43	500
macro avg	0.56	0.43	0.42	500
weighted avg	0.56	0.43	0.42	500



```
1 print('Accuracy:',accuracy_score(y_trues, y_preds))
2 print('Precision:',precision_score(y_trues, y_preds, average='macro'))
3 print('Recall:',recall_score(y_trues, y_preds, average='macro'))
4 print('f1-Score:',f1_score(y_trues, y_preds, average='macro'))
```


```
Accuracy: 0.428
Precision: 0.5615894618175594
Recall: 0.42800000000000005
f1-Score: 0.4175131160752478
```

```
1 plot_predict(model1)
```


Prediction: 3



Prediction: 9




Prediction: 4



Prediction: 9



Prediction: 9



Prediction: 0

